

## ADL HW4 Report

### Explain the structure of your networks and loss terms in detail (1.5%)

Basically I used the structure similar in DCGAN for both the generator and discriminator. I have added a conv layer to both the generator and discriminator to resize the output of the image to 128x128. I have also taken off the sigmoid function in the discriminator layer and batch normalization layers in the discriminator for WGAN-GP training and add 4 classifiers(Hair classifier, eye classifier, face classifier, glass classifier) to classify the conditions.

Generator:

```
Generator(  
    (gen): Sequential(  
        (0): ConvTranspose2d(115, 1024, kernel_size=(4, 4), stride=(1, 1), bias=False)  
        (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        (2): ReLU(inplace)  
        (3): ConvTranspose2d(1024, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
        (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        (5): ReLU(inplace)  
        (6): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
        (7): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        (8): ReLU(inplace)  
        (9): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
        (10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        (11): ReLU(inplace)  
        (12): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
        (13): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        (14): ReLU(inplace)  
        (15): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
        (16): Tanh()  
    )  
)
```

**Discriminator:**

```

Discriminator(
    (conv_layers): Sequential(
        (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (1): LeakyReLU(negative_slope=0.2, inplace)
        (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (3): LeakyReLU(negative_slope=0.2, inplace)
        (4): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (5): LeakyReLU(negative_slope=0.2, inplace)
        (6): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (7): LeakyReLU(negative_slope=0.2, inplace)
        (8): Conv2d(512, 1024, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (9): LeakyReLU(negative_slope=0.2, inplace)
    )
    (discriminator_layer): Sequential(
        (0): Conv2d(1024, 1, kernel_size=(4, 4), stride=(1, 1))
    )
    (bottleneck): Sequential(
        (0): Conv2d(1024, 1024, kernel_size=(4, 4), stride=(1, 1))
        (1): LeakyReLU(negative_slope=0.2)
    )
    (hair_classifier_layer): Sequential(
        (0): Linear(in_features=1024, out_features=128, bias=True)
        (1): LeakyReLU(negative_slope=0.2)
        (2): Linear(in_features=128, out_features=6, bias=True)
        (3): LogSoftmax()
    )
    (eye_classifier_layer): Sequential(
        (0): Linear(in_features=1024, out_features=128, bias=True)
        (1): LeakyReLU(negative_slope=0.2)
        (2): Linear(in_features=128, out_features=4, bias=True)
        (3): LogSoftmax()
    )
    (face_classifier_layer): Sequential(
        (0): Linear(in_features=1024, out_features=128, bias=True)
        (1): LeakyReLU(negative_slope=0.2)
        (2): Linear(in_features=128, out_features=3, bias=True)
        (3): LogSoftmax()
    )
    (glass_classifier_layer): Sequential(
        (0): Linear(in_features=1024, out_features=128, bias=True)
        (1): LeakyReLU(negative_slope=0.2)
        (2): Linear(in_features=128, out_features=2, bias=True)
        (3): LogSoftmax()
    )
)
)

```

**Loss terms:****Adversarial loss:**

GAN	DISCRIMINATOR LOSS	GENERATOR LOSS
WGAN GP	$\mathcal{L}_D^{WGANGP} = \mathcal{L}_D^{WGAN} + \lambda \mathbb{E}_{\hat{x} \sim p_g} [(\ \nabla D(\alpha x + (1 - \alpha \hat{x}))\ _2 - 1)^2]$	$\mathcal{L}_G^{WGANGP} = -\mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})]$

$$\lambda = 10$$

**Classifier loss:**

Generator classifier loss
$CrossEntropyLoss(P(class = hair   X_{fake}), fake\_tag_{hair})$
$+ CrossEntropyLoss(P(class = eye   X_{fake}), fake\_tag_{eye})$
$+ CrossEntropyLoss(P(class = face   X_{fake}), fake\_tag_{face})$
$+ CrossEntropyLoss(P(class = glasses   X_{fake}), fake\_tag_{glasses})$

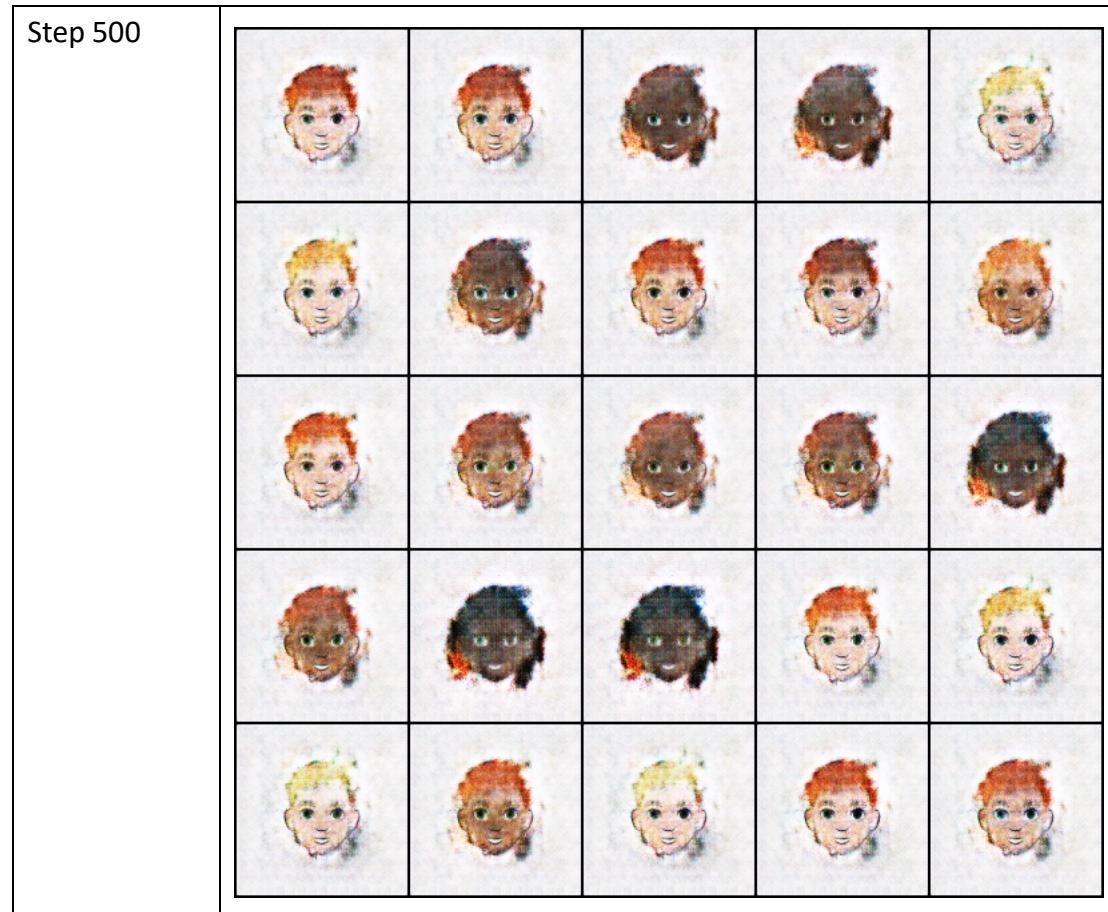
**Discriminator classifier loss**

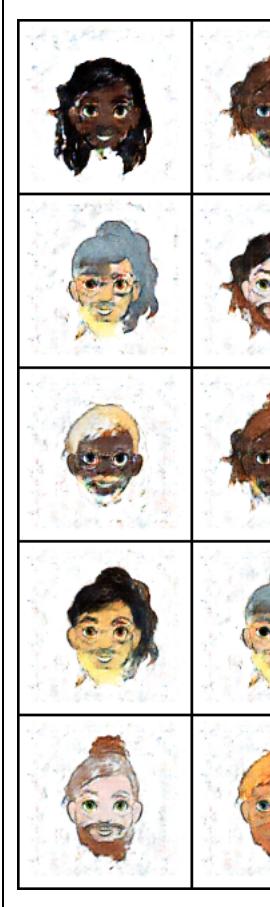
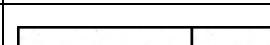
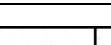
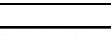
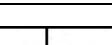
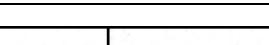
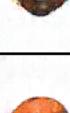
$$\begin{aligned} & \text{CrossEntropyLoss}(P(\text{class} = \text{hair}|X_{\text{real}}), \text{real\_tag}_{\text{hair}}) \\ & + \text{CrossEntropyLoss}(P(\text{class} = \text{eye}|X_{\text{real}}), \text{real\_tag}_{\text{eye}}) \\ & + \text{CrossEntropyLoss}(P(\text{class} = \text{face}|X_{\text{real}}), \text{real\_tag}_{\text{face}}) \\ & + \text{CrossEntropyLoss}(P(\text{class} = \text{glasses}|X_{\text{real}}), \text{real\_tag}_{\text{glasses}}) \end{aligned}$$

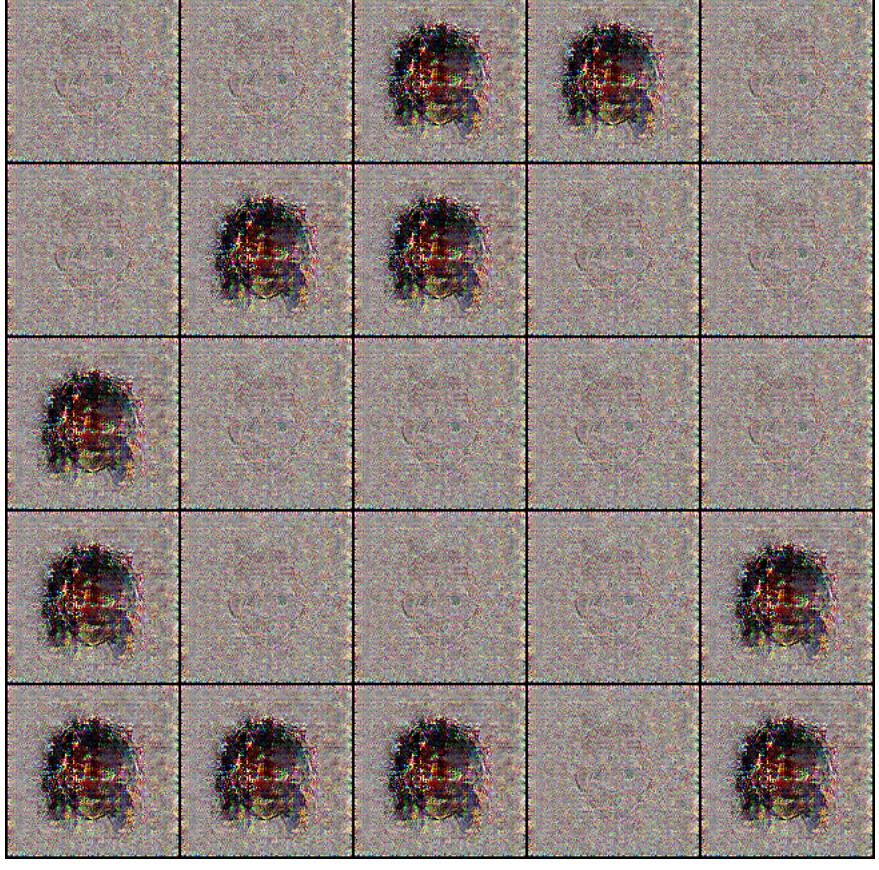
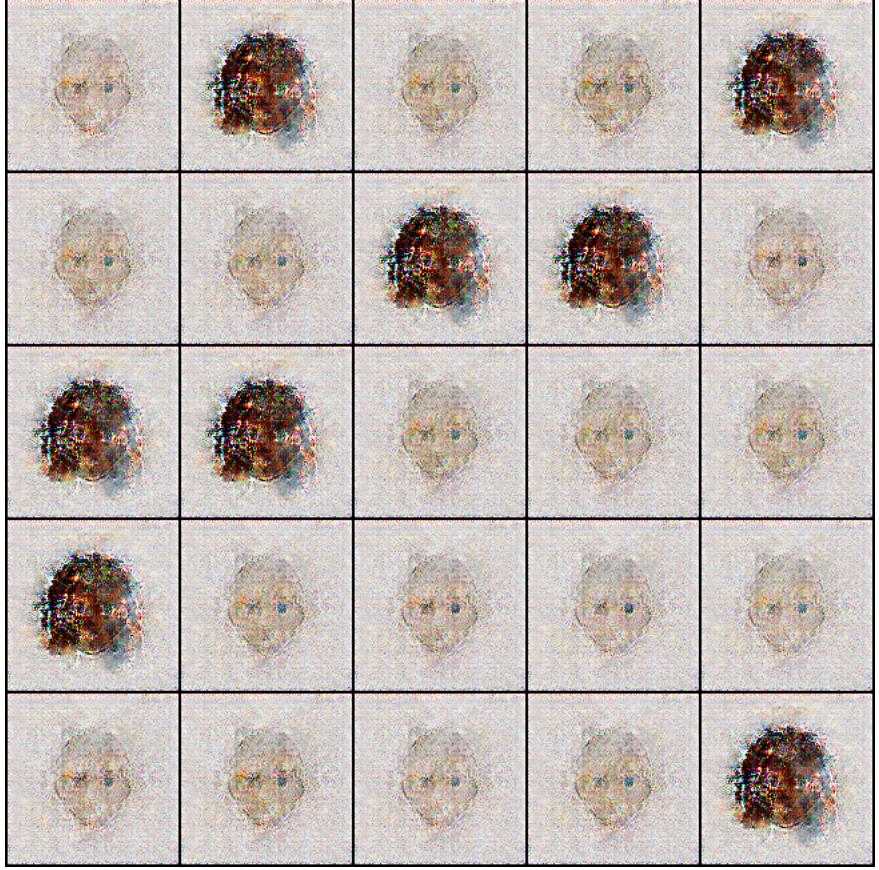
**Experiment settings and results**

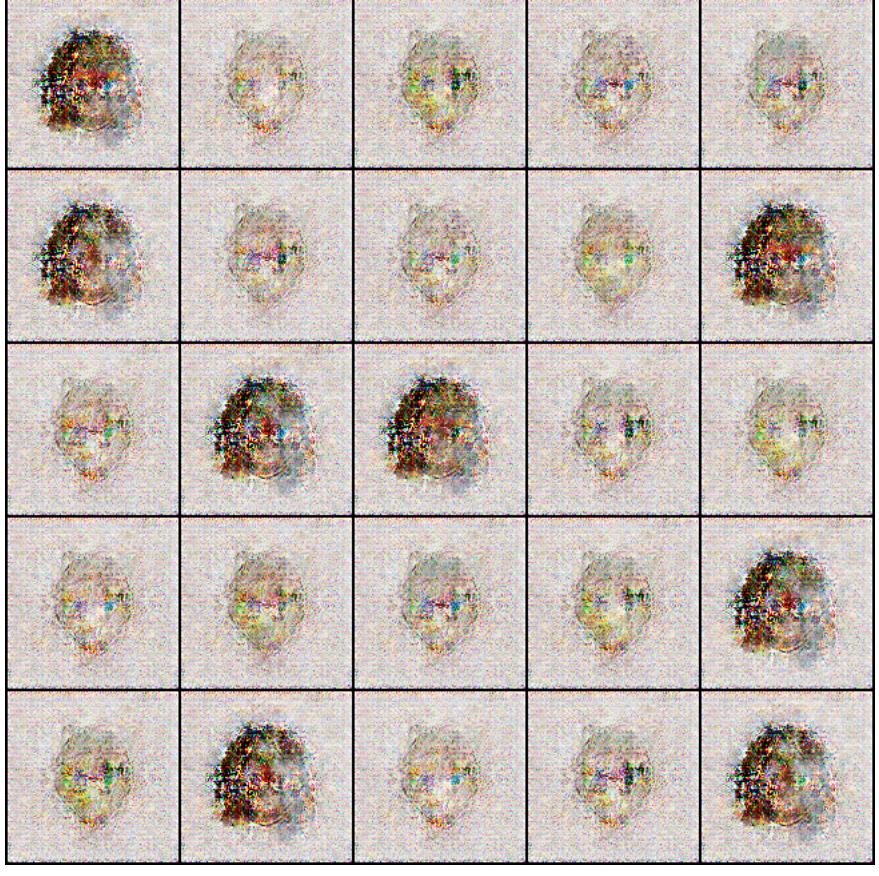
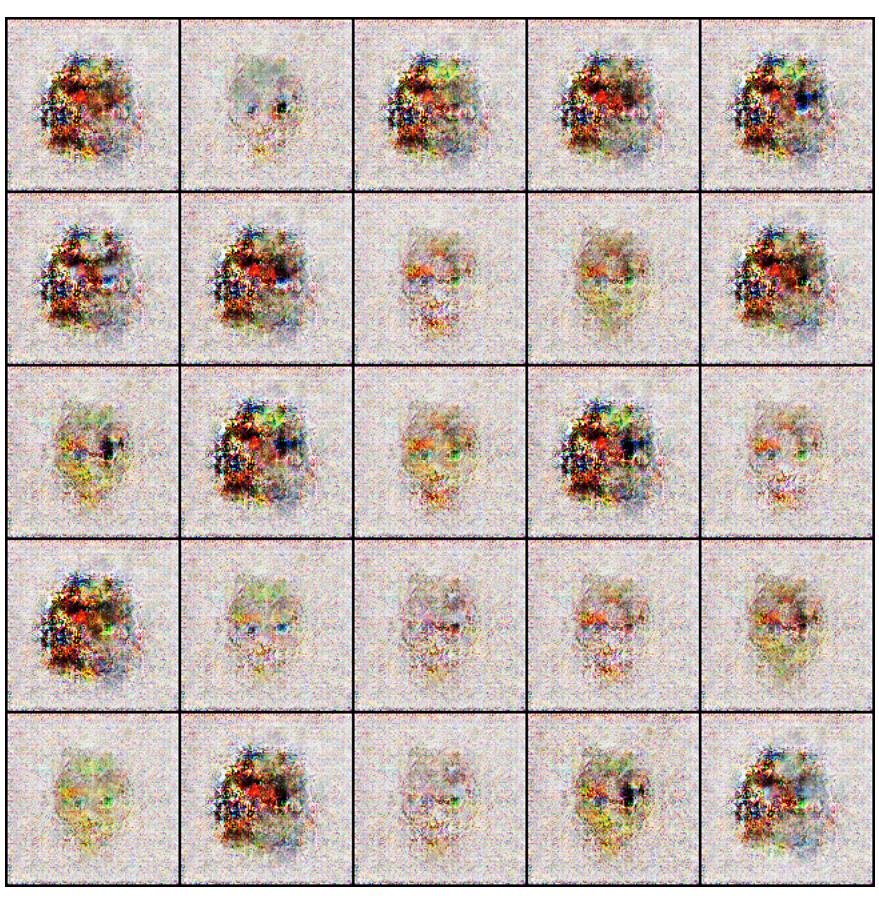
Iter	50000
Discriminator_iter: Generator_iter when training	5:1
Batch size	64
Learning Rate	0.0001
Optimizer	Adam
Latent dim	100
FID score	17.277

**Plot your training progress (10 pics) (0.5%)**



Step 5000								
Step 10000								
								
								
								
								
								
Step 20000								
								
								
								
								

Step 25000	
Step 30000	

Step 35000	
Step 40000	

Step 45000					
Step 50000					

可以觀察到訓練到一半時，generator一度有崩潰的情形。但是後來有比較好一點。

**Design at least 3 different experiments. Describe your settings, making comparisons and report your observations. (4% + 0.5% bonus)**

$X_{fake} = G(z, tag)$  is the fake image(or generated image),  
and  $X_{real}$  is the real image

1.

Discriminator Loss function
$-\left(\log P(S = \text{real} X_{real}) + \log\left(1 - P(S = \text{fake} X_{fake})\right)\right) * \frac{1}{2}$
$-(\log P(\text{class} = c X_{real}))$
Generator Loss function
$-\log P(S = \text{fake} X_{fake}) - \log P(\text{class} = c X_{fake})$

The classification loss for fake images are omitted, since we believe that fake images will confuse the discriminator. I divide the adversarial loss for discriminator by 2 to balance the gradient descent schedule of G/D.

Iter	Batch size	FID score
50000	128	187.534

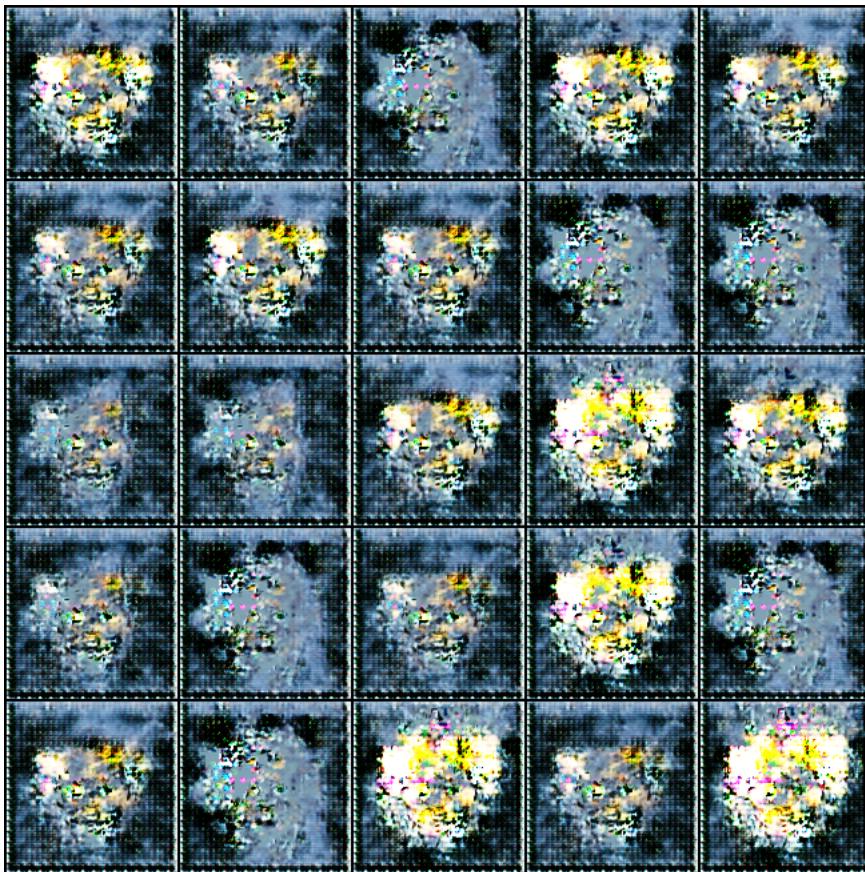
2.

Discriminator Loss function
$-\left(\log P(S = \text{real} X_{real}) + \log\left(1 - P(S = \text{fake} X_{fake})\right)\right) * \frac{1}{2}$
$-(\log P(\text{class} = c X_{real}) + \log P(\text{class} = c X_{fake}))$
Generator Loss function
$-\log P(S = \text{fake} X_{fake}) - \log P(\text{class} = c X_{fake})$

Add classification loss for fake images in the discriminator loss function.

Iter	Batch size	FID score
300000	32	Generator crashed

訓練到約 25 萬 iters 時 generator 崩潰(見下圖)。所以終止訓練。



3.

Discriminator Loss function

$$\begin{aligned}
 & - \left( \log P(S = \text{real} | X_{\text{real}}) + \log \left( 1 - P(S = \text{fake} | X_{\text{fake}}) \right) \right) \\
 & - (\log P(\text{class} = \text{hair} | X_{\text{real}}) + \log P(\text{class} = \text{eye} | X_{\text{real}}) \\
 & \quad + \log P(\text{class} = \text{face} | X_{\text{real}}) + \log P(\text{class} = \text{glass} | X_{\text{real}}))
 \end{aligned}$$

Generator Loss function

$$\begin{aligned}
 & - \log P(S = \text{fake} | X_{\text{fake}}) - \log P(\text{class} = \text{hair} | X_{\text{fake}}) \\
 & \quad - \log P(\text{class} = \text{eye} | X_{\text{fake}}) - \log P(\text{class} = \text{face} | X_{\text{fake}}) \\
 & \quad - \log P(\text{class} = \text{glass} | X_{\text{fake}})
 \end{aligned}$$

Split the classifier in discriminator into 4 different parts to classify hair colors, eye colors, face colors and with and without glasses.

Iter	Batch size	FID score
50000	128	133.792